

EUNKYU (EUNICE) PARK

(+1) 412-583-5837 | eunkyup@andrew.cmu.edu | eunkyu.park@vision.snu.ac.kr | <https://dbsltm.github.io/>

EDUCATION

Seoul National University

Integrated M.S.-Ph.D. in Artificial Intelligence

Seoul, South Korea

March 2021 — Present

- Advised by Professor Gunhee Kim (VISION & LEARNING LAB)
- Research on **how models reason and how people come to trust them**: (1) evaluating the trustworthiness of large language models (faithfulness of reasoning, localizing perception-to-understanding failures); (2) studying how people perceive, rely on, and are misled by AI explanations; and (3) building benchmarks and interfaces to better evaluate and capture social reasoning. I design human-subject studies to evaluate how users perceive and calibrate trust in these systems and translate those findings into better learning signals for model alignment.

Columbia University in The City of New York, *Fu Foundation School of Engineering and Applied Science* New York, NY

B.S. in Computer Science, Intelligent Systems

Aug 2017 — May 2020

Bard College

B.A. in Mathematics

Great Barrington, MA

Aug 2014 — May 2020

PUBLICATIONS (*denotes equal contribution)

MM-SCALE: Evaluating Evidence-Grounded Moral Judgment in Vision-Language Models *NeurIPS 2026 E&D*

Eunkyu Park, Wesley Hanwen Deng, Cheyon Jin, Matheus Kunzler Maldaner, Jordan Wheeler, Jason I. Hong, Hong Shen, Adam Perer, Ken Holstein, Motahhare Eslami, Gunhee Kim (with CMU)

Critical or Compliant? The Double-Edged Sword of Reasoning in Chain-of-Thought Explanations *EMNLP 2026*

Eunkyu Park, Wesley Hanwen Deng, Vasudha Varadarajan, Mingxi Yan, Gunhee Kim, Maarten Sap*, Motahhare Eslami* (with CMU)

Jury Duty: Calibration and Orientation Failures in MLLM-as-a-Judge Under Cultural Ambiguity *EMNLP 2026*

Daniel Lee, Harshit Sharma, Eunkyu Park, Pranav Narayanan Venkit, Jeonghwan Kim, Kah Mun Chia, Andreas Vlachos, Shafiq Joty (with Salesforce)

Cognitive Chain-of-Thought (CoCoT): Structured Multimodal Reasoning about Social Situations *COLM 2026*

Eunkyu Park, Wesley Hanwen Deng, Gunhee Kim, Motahhare Eslami, Maarten Sap (with CMU)

HalLoc: Token-level Localization of Hallucinations for Vision Language Models

CVPR 2025

Eunkyu Park*, Minyeong Kim*, Gunhee Kim

Who Does What in AI Auditing? Designing Human-AI Collaboration for Auditing Generative AI *arXiv 2026*

Eunkyu Park, Markelle Roesti, Wesley Hanwen Deng, Renata Barreto, Mohammad Tahaei, Ken Holstein, Jason I. Hong, Motahhare Eslami (with eBay, CMU)

EXPERIENCE

Carnegie Mellon University

Visiting Researcher

Pittsburgh, PA

June 2025 — Present

With Motahhare Eslami (HCII) and Maarten Sap (LTI)

- Run human-subjects experiments and user studies on how people trust and rely on AI reasoning, including an 800-participant online experiment showing that people scrutinize model reasoning less, and miss more of its errors, when they agree with its conclusion

- Designed a between-subjects study that separates whether an AI advisor's decision or its reasoning matches each participant's values, measuring effects on trust and willingness to delegate
- Built interfaces for human-AI collaboration, including a platform where practitioners and AI work together to audit AI agents, with a practitioner study, in collaboration with eBay
- Developed a benchmark of about 32K multimodal scenarios with human moral acceptability ratings to evaluate whether vision-language models ground their moral judgments in the right evidence
- Developed reasoning frameworks for moral reasoning in vision-language models, with probing analyses of how each stage is represented inside the model

VISION & LEARNING LAB at Seoul National University

Seoul, South Korea

Research Intern

July 2020 — March 2021

- Assisted research in developing models for Video Question and Answering benchmark that could expand to Drama-QA, TV-QA domains, eventually expanding to a multi-modal commonsense understanding framework
- Ranked top-10% in the 2020 AI Challenge hosted by the Ministry of Science and Technology

PION CORPORATION

Seoul, South Korea

Machine Learning Engineer Intern

July 2019 — Sep 2019, April 2020 — July 2020

- Developed a model extracting information such as an object's position, label, colors, coordinates from videos/images by masking objects with R-CNN and identifying them with YOLOv3 in a single frame
- Used Selenium for test-running the application, taking input images/videos from client to extract metadata and feed into the GAN training model

DATA SCIENCE INSTITUTE at Columbia University

New York, NY

Undergraduate Research Assistant

Feb 2020 — May 2020

- Analyzed text data to investigate the relationship between users' demographic information and sentiment towards self-driving cars
- Studied patterns of interactions among users using python and visualized using Tableau
- Labeled comments with sentiment score to train a classification model to cluster segments by relevant topic.

ACADEMIC SERVICE

- **Reviewer:** AAAI 2027, COLM 2026, IJCV 2025
- **Teaching Assistant at Seoul National University:** Computer Vision, Probabilistic Graphical Models, General Artificial Intelligence
- **Teaching Assistant for Corporate Training:** 2025 LG Energy Deep Learning Course

QUALIFICATIONS

- **Human-subjects research:** experimental and survey design (between-subjects experiments, user studies, practitioner interviews), large-scale online studies on Prolific and Amazon Mechanical Turk
- **Study platforms:** web-based study interfaces and experiment instruments (HTML/JavaScript), LLM-based user simulation for piloting study designs
- **Analysis:** Python, statistical modeling
- **Machine learning:** PyTorch, TensorFlow
- Fluent in English, Korean
- Interests in golf, playing the violin